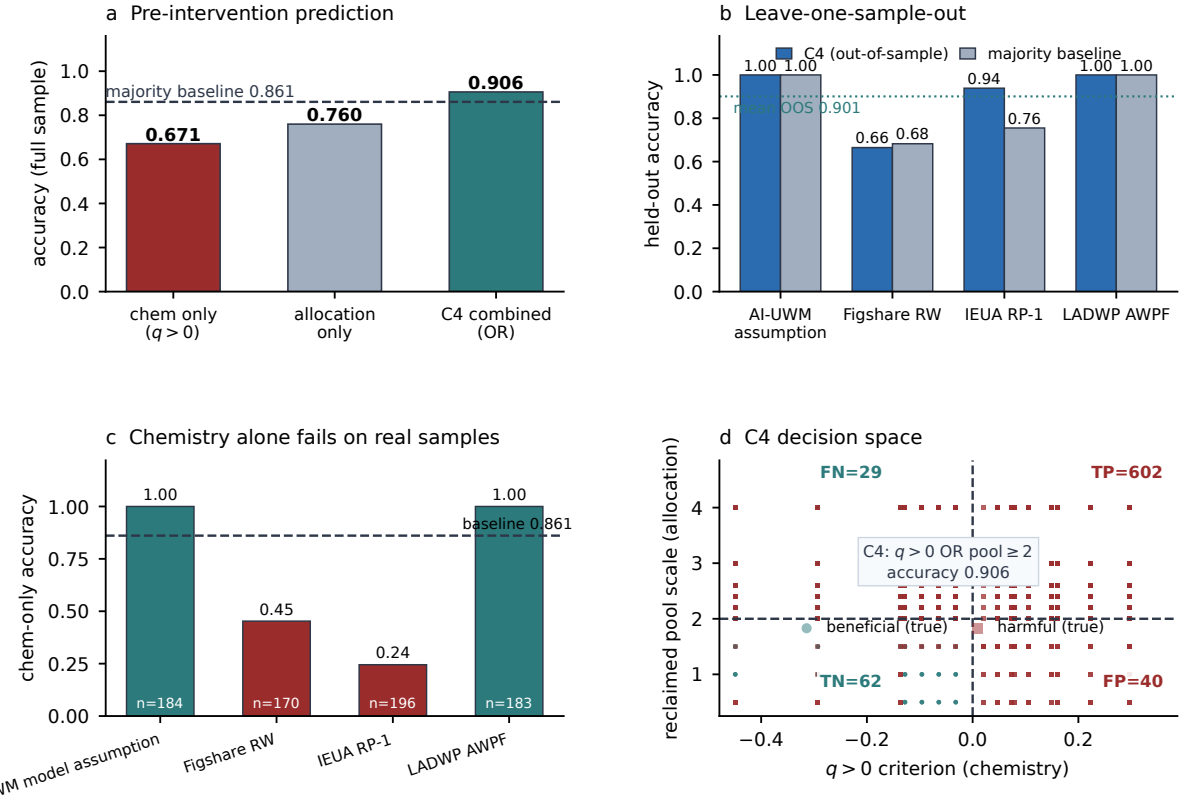

Data: A31_WP2/wp2_e3_independent_prediction_summary.json, wp2_e3b_combined_prediction_summary.json, wp2_e3_predictions.csv, wp2_e3b_features.csv.
n_all = 813 reliable scenarios; n_eff = 733 after excluding 80 near-zero rows (|dS_pot| < 1e-9). Leave-one-sample-out groups: n=184/170/196/183.
LIMITATION: the predictor is causal only for the MECHANISM it encodes (pre-intervention chemistry and pool availability); it requires allocation-side data and gives no city-level calibration. AI-UWM and Figshare RW are incomplete-species samples (2-3 of 5 species), so their per-sample accuracies are SI-only.